Research and Build

- Product Thesis
  - We are building a scalable technology platform for conducting high-stakes online exams (e.g., CAT 2025 and beyond) with a projected volume of **100k–2 million aspirants per test**.
  - Each test will:
    - Consist of **50-300 MCQs**
    - Have each question tagged to **one of 19 possible subjects**
    - Be time-bound: **3-hour duration**
    - Have a taker volume ranging from **10,000 to 1,000,000+ aspirants per test**
    - Be used across **100K–1M such test instances annually**
  - Your Objective
    - You are expected to design and prototype the **backend foundation** for this platform with a focus on **performance, accuracy, and scalability**. This involves tackling three core problems.
  - Problem 1: Efficiently Storing User Submissions
    - Design a schema and submission workflow that:
      - Stores answers submitted by each user for each question
      - Tracks the correctness of the answer (correct: +5, wrong: -1)
      - Supports **subject-wise scoring** for each test taker
      - Can scale to tens of millions of records without bottlenecks
    - *Key Expectations:*
      - Document schema design (MongoDB collections)
      - FastAPI routes like `POST /submit-answer`
      - Indexing strategies for performance
      - Efficient querying for result generation
  - Problem 2: Publishing Final Test Results (Ranks & Percentiles)
    - Once the exam concludes, the system must:
      - Calculate **overall score** and **subject-wise scores** for all users
      - Compute:
        - **Ranks**
        - **Percentiles**
        - **Subject-wise percentiles**
      - Efficiently publish the final leaderboard for 1M+ users
    - *Key Expectations:*

        - Algorithm/logic to compute percentiles & ranks at scale
        - Trade-offs between pre-computation vs on-demand generation
        - Thought process around caching, async jobs, data structure choices
- Problem 3: Predictive Results for Future Mock Test Takers
    - After the real exam ends and the result is published:
        - Future aspirants (e.g. CAT 2025 aspirants) will take **mock tests** modeled after CAT 2024
        - For each mock test user, the system must provide:
            - **Predicted Rank**
            - **Predicted Percentile**
            - **Predicted Subject-wise Percentile**
    - 💭 *What we're looking for:*
        - Approach to compare mock user's score with real test result distribution
        - Insight into how to model predictions using existing data
        - Thoughts on real-time prediction performance vs accuracy tradeoffs
- Development Expectations
    - You can use **FastAPI** (or any equivalent Python framework) to:
        - Expose REST APIs
        - Use an MongoDB DataBase
        - Show modular, readable, and scalable code structure
- Bonus (Not Mandatory but Appreciated)
    - Rate limiting or anti-cheat mechanisms
    - Basic auth or user ID simulation